

# Computer modeling of physical phenomena






Lecture VII: Game theory

tomasz.szawello@fuw.edu.pl

fuw.edu.pl/~tszawello/cmpp2023

room 4.45


consultation hours: Mondays 11.15-13.00




# Game theory

Game theory studies the ways in which strategic interactions among rational players produce outcomes with respect to the players' preferences (the outcomes might not have been intended by any of them).

# First book on game theory



Since war is a game...



Sun Tzu, 544–496 BC



# Modern game theory




Neumann, J., and Morgenstern, O., (1947). The Theory of Games and Economic Behavior.



Nash, J., (1950). Equilibrium Points in n-Person Games, PNAS, 36, 48-49.

Nash, J., (1950). Non-Cooperative Games, PhD Thesis, Princeton.

# John and Alicia Nash



A Beautiful Mind (2001)

# A Beautiful Mind...



# Rules, Strategies, Payoffs, and Equilibrium

- The **rules** of the game state who can do what, and when they can do it.
- A player's **strategy** is a plan for actions in each possible situation in the game.
- A player's **payoff** is the amount that the player wins or loses in a particular situation in a game.
- A player has a **dominant strategy** if his best strategy doesn't depend on what other players do.



# Zero-sum or not


**Zero-sum** describes a situation in which a participant's gain (or loss) is exactly balanced by the losses (or gains) of the other participant(s).



Non-zero-sum games are more complex to analyze **and more interesting!**

# Nash Equilibrium


A situation when each player's strategy is optimal, given the strategies of the other players (no player can benefit from unilaterally changing his strategy, while all other players stay fixed).




Every finite game has at least one Nash equilibrium.

# Nash Equilibrium

HEY, DR. NASH, I THINK THOSE GALS OVERTHERE ARE EYEING US. THIS IS LIKE YOUR NASH EQUILIBRIUM, RIGHT? ONE OF THEM IS HOT, BUT WE SHOULD EACH FLIRT WITH ONE OF HER LESS-DESIRABLE FRIENDS. OTHERWISE WE RISK COMING ON TOO STRONG TO THE HOT ONE AND JUST DRIVING THE GROUP OFF.



WELL, THAT'S NOT REALLY THE SORT OF SITUATION I WROTE ABOUT. ONCE WE'RE WITH THE UGLY ONES, THERE'S NO INCENTIVE FOR ONE OF US NOT TO TRY TO SWITCH TO THE HOT ONE. IT'S NOT A STABLE EQUILIBRIUM.



CRAP, FORGET IT. LOOKS LIKE ALL THREE ARE LEAVING WITH ONE GUY.



## Prisoner's Dilemma



Merrill Flood and Melvin Dresher (RAND, 1950)



# Intro: Buick sale



In June 1949, Flood wanted to buy a used Buick from a RAND employee who was moving back East. Buyer and seller were friends. They weren't looking to cheat each other, just to agree on a fair price for the car. How should they set a price?

Dealer: buy for 500\$, sell for 800\$, should they split an extra \$300 profit?

What if one defects (I will sell for 750\$!) - the other is still better off to accept it than to go the dealer.


**The party who is more unreasonable is apt to get the better of the deal!**

based on: William Poundstone, *Prisoner's Dilemma*, 1992

# Prisoner's dilemma

Bonnie and Clyde are arrested by the police and charged with various crimes. They are questioned in separate cells, unable to communicate with each other. They know how it works:

- if both resist interrogation (cooperating with each other) and proclaim mutual innocence, they will get a 2-year sentence for robbery,
- if one confesses (defecting) and the other doesn't (cooperating), the confesser goes free and the other will get a severe 8-year sentence,
- if they both confess (defecting), then the judge will sentence both to moderate 4 years in prison.




What should Bonnie do? What should Clyde do?

# The payoff matrix


		CLYDE	
		Confess	Not confess
BONNIE	Confess	4 years for B 4 year for C	0 years for B 8 years for C
	Not confess	8 years for B 0 years for C	2 years for B 2 years for C

# Bonnie's Decision Tree


**If Clyde confesses**




Best strategy



**If Clyde does not confess**




Best strategy



The dominant strategy for Bonnie is to confess (defect) because no matter what Clyde does, she is better off confessing.




It seems we should always defect and never cooperate...



# Examples...

- Cheating on a cartel:
  - Cartel members can charge the monopoly price or a lower price.
  - The best strategy is to charge the low price.
- Trade wars between countries
  - Free trade benefits both trading countries.
  - Tariffs can benefit one trading country.
  - Imposing tariffs can be a dominant strategy and establish a Nash equilibrium even though it may be inefficient.
- Advertising
  - All firms advertising tends to equalize the effects.
  - Everyone would gain if no one advertised.

# Examples...



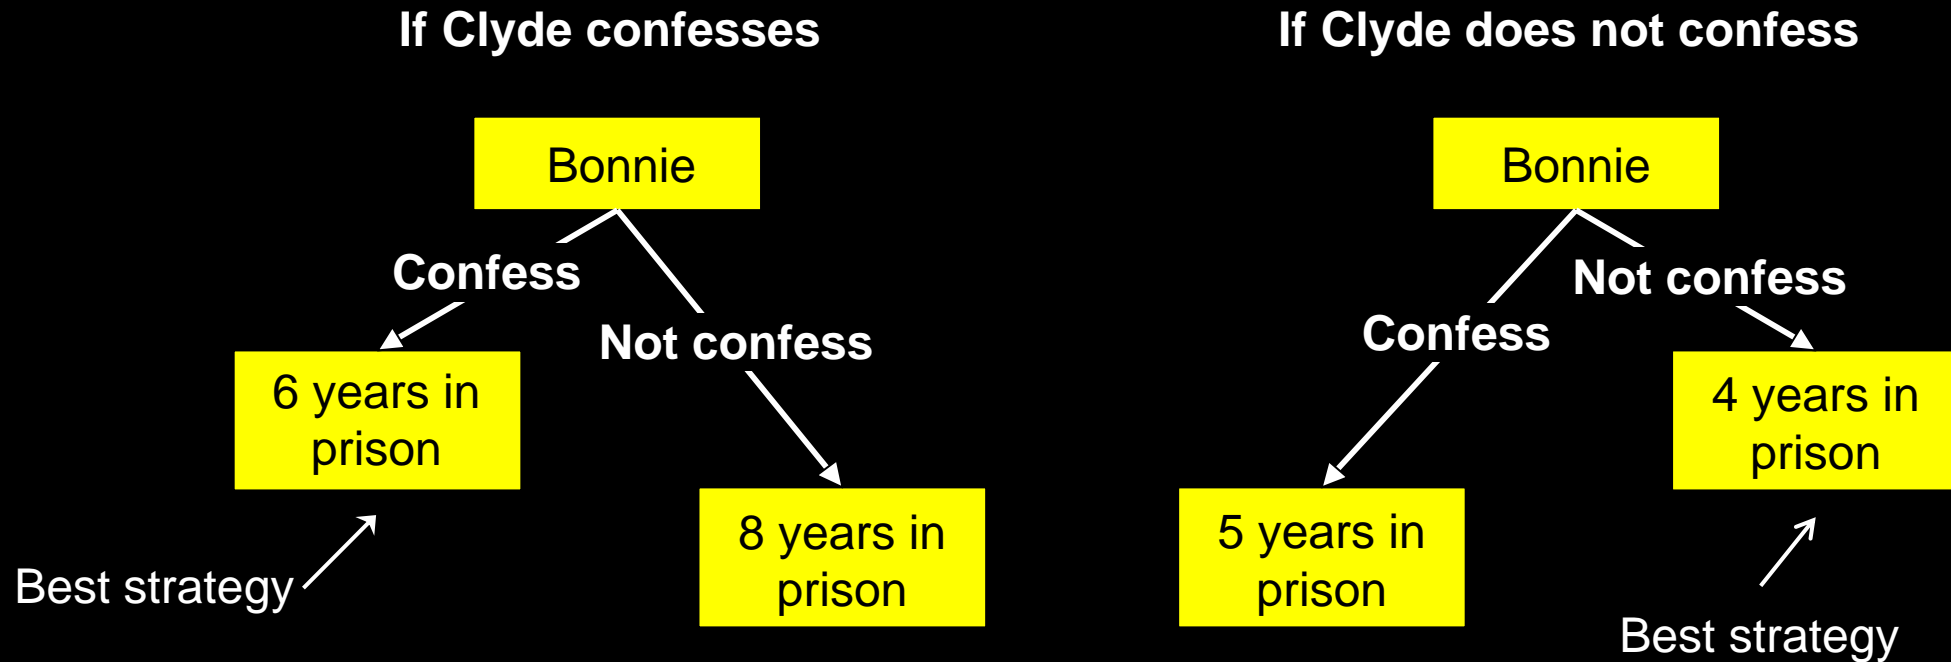
# Games without dominant strategies

In many games the players have no dominant strategy – often a player's strategy depends on the strategies of others.

		CLYDE	
		Confess	Not confess
BONNIE	Confess	6 years for B 1 year for C	5 years for B 3 years for C
	Not confess	8 years for B 0 years for C	4 years for B 2 years for C

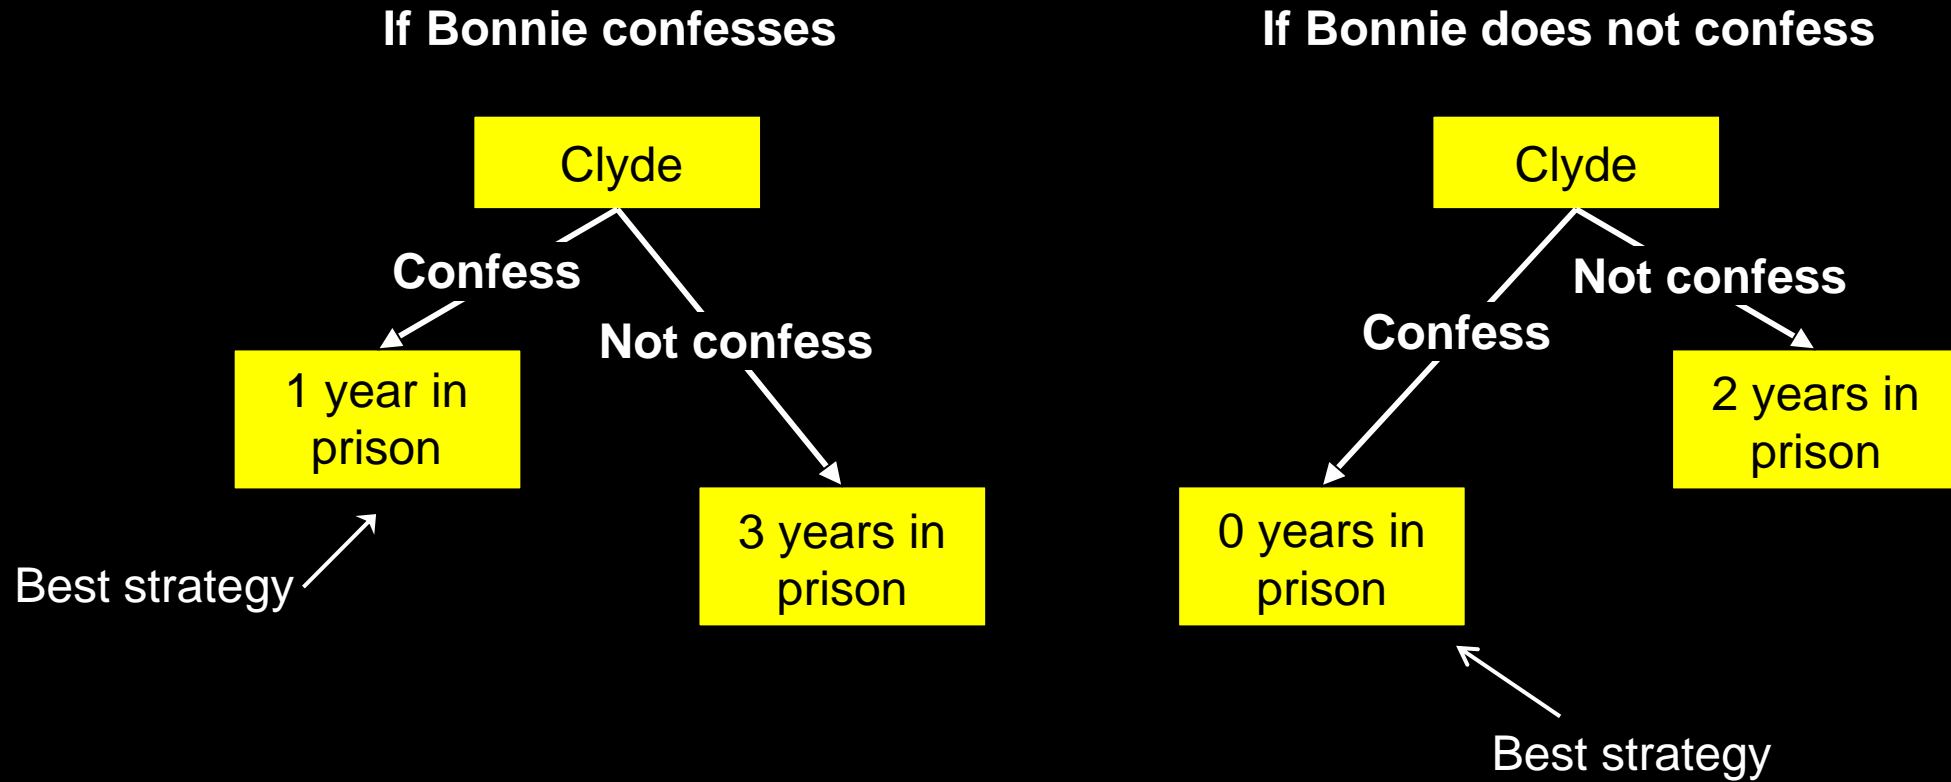


# Bonnie's Decision Tree



Bonnie has no explicit dominant strategy, but there is an *implicit* one!

# Clyde's Decision Tree



Clyde has a dominant strategy (confess); thus, Bonnie could assume that he would choose it and choose to confess as well.

## Some games have no simple solution


In *matching pennies*, neither player has a dominant strategy:

		Player B	
		Heads	Tails
Player A	Heads	1, -1	-1, 1
	Tails	-1, 1	1, -1

This game has no pure strategy Nash equilibrium since there is no pure strategy (heads or tails) that is a best response to a best response.

# Mixed (non-deterministic) strategy

How should the players play, possibly in a randomized way, to maximize their expected reward?



**Randomized** Nash equilibrium!



# Randomized Nash equilibrium

These best response functions have only one intersection:  $(0.5, 0.5)$ . Thus there is only one pair of choices for  $(P_A, P_B)$  where no player has an incentive to deviate unilaterally. If, e.g.  $(P_A, P_B) = (0.4, 0.5)$ , then B has an incentive to deviate from 0.5 because changing  $P_B$  to 1 will increase his chance of winning. Only the intersection of the two best response functions is such that no player wants to deviate unilaterally: this is Nash Equilibrium. Thus, the game of Matching Pennies has a unique Nash Equilibrium, and it is randomized.

# Payoff matrix for the generic symmetric two-person dilemma game



(A's payoff, B's payoff)

		Player B	
		cooperate	defect
Player A	cooperate	(CC,CC) reward for mutual cooperation	(CD,DC) sucker's payoff and temptation to defect
	defect	(DC,CD) temptation to defect and sucker's payoff	(DD,DD) punishment for mutual defection

# Payoffs


Four payoffs involved:

- CC: both players cooperate,
- CD: you cooperate but other defects (aka "sucker's payoff"),
- DC: you defect and other cooperates (aka "temptation to defect"),
- DD: both players defect.

Some classical dilemma games:

- Prisoner's dilemma:  $DC > CC > DD > CD$
- Chicken:  $DC > CC > CD > DD$
- Stag Hunt / Trust dilemma:  $CC > DC > DD > CD$

# Chicken




Rebel Without a Cause (1955)

# Chicken

$DC > CC > CD > DD$

Cooperation: swerving

Defecting: not swerving



- The worst thing that can happen is for both players *not* to swerve.
- The best thing that can happen, is to show your machismo by not swerving and letting the other driver swerve.
- Being chicken is the next to worst outcome, but still better than dying.

There is a cooperative outcome in chicken. It's not so bad if both players swerve. Both come out alive, and no one can call the other a chicken.



# Example payoff matrix

		Player B	
		swerve	drive straight
Player A	swerve	2, 2	1, 3
	drive straight	3, 1	0, 0

- No dominant strategy: the best is to do the opposite to what the other player does.
- Playing safe? Swerving has the maximum minimum.
- **Chicken vs the prisoner's dilemma:** mutual defection (the crash when both players drive straight) is the most feared outcome in chicken.

# Trust Dilemma

CC > DC > DD > CD

		Player B	
		haircut	no haircut
Player A	haircut	5, 5	0, 4
	no haircut	4, 0	2, 2



- The best possible outcome would be for both of you to get the haircut.
- It would be awful to be the *only* one to show up with the haircut.
- If you *didn't* get the haircut, but the friend did, and looked like a real jerk, that would almost be as good as if both of you got the haircut.
- It wouldn't really be so bad if *no one* got the haircut.

# Trust Dilemma

$CC > DC > DD > CD$

		Player B	
		haircut	no haircut
Player A	haircut	5, 5	0, 4
	no haircut	4, 0	2, 2



**risk dominant strategy** - no haircut to minimize the risk

**payoff dominant strategy** - haircut to get the maximum reward

The more uncertainty players have about the actions of the other players, the more likely they will choose the risk dominant strategy.

# More examples of the GT in real life

## **Communal coffeepot**

- Cooperate by making a new pot of coffee if you take the last cup.
- Defect by taking the last cup and not making a new pot, depending on the next coffee seeker to do it.
- $DC > CC > DD > CD$

## **Class team project**

- Cooperate by doing your part well and on time.
- Defect by slacking, hoping the other team members will come through and sharing the benefit of a good grade.
- $DC > CC > DD > CD$

# How to save morality?



# Iterated Prisoner's Dilemma

Game theory shows that a rational player should always defect when engaged in a prisoner's dilemma situation. We know that in real situations, people don't always do this.

Why not? Possible explanations:

- people aren't rational,
- morality,
- social pressure,
- fear of consequences,
- evolution of species-favoring genes.

In frames of game theory, some of these ideas can be explored by playing the iterated version of the games.



# Iterated Prisoner's Dilemma

**Key idea:** in many situations, we play more than one “game” with a given player.

Players have complete knowledge of the past games, including their choices and the other player's choices. Your choice in future games when playing against a given player can be partially based on whether he has been cooperative in the past.



# Axelrod & Hamilton tournament

Round-robin computer tournament between strategies submitted by colleagues.

Each game between competing strategies consisted of 200 moves, yielding possible scores from 0 to 1000, so that a "good performance" score was judged to be 600 points (complete mutual cooperation) and a "poor performance" score would be 200 points (complete mutual defection).

The overall winner with an average of 504 points per game, submitted by game theorist Anatol Rapoport, was so called Tit-for-Tat: a player cooperates on the first move and then does exactly what the other player did on the previous move.

While never actually attaining a higher score in any particular game, this strategy can lead to many rounds of cooperation between opponents with its commensurate rewards, whereas the more defecting strategies do poorly.

# Axelrod on Tit-for-Tat

"What accounts for Tit-for-Tat's robust success is its combination of being nice [not being the first to defect] , retaliatory, forgiving, and clear. Its niceness prevents it from getting into unnecessary trouble. Its retaliation discourages the other side from persisting whenever defection is tried. Its forgiveness helps restore mutual cooperation. And its clarity makes it intelligible to the other player, thereby eliciting long-term cooperation" (Axelrod [1984]).




# GA for Prisoner's Dilemma

You can try to do better than TfT by employing genetic algorithms.

Some hints can be found here:

<http://www-personal.umich.edu/~axe/research/Evolving.pdf>




# Spatial version of PD

Cooperation is easier to maintain in a sedentary population: defectors can thrive in an anonymous crowd, but mutual aid is frequent among neighbors.

As a model of the above, one can consider PD played on a lattice with either Moore or von Neumann neighbourhood.

- Various strategies are allowed to compete against each other.
- Initially, strategies are randomly assigned to the cells.
- Each cell then competes for a fixed number of rounds with each of its neighbors.
- At the end of each session a cell will adopt the strategy of its most successful neighbor in terms of highest total score.

# Spatial version of PD



The evolution of the spatial iterated prisoner's dilemma with five competing strategies from an initial random configuration. Each cell's strategy plays off five times against its eight neighbors and then adopts the strategy of its most successful neighbor. Images are taken at the initial state, after 2, 5, 10, and 50 iterations.




Pavlov - repeats its former choice whenever it wins a round (C/C or D/C) and switches otherwise (if winning, hold fast; if not, change course).

TfT seems to win, but this depends on initial conditions!



# The same with noise



A certain level of noise introduced, that is, a cell will make a random move in a single round, thus increasing the dynamics of the evolution.



Can you devise a better strategy for spatial PD using GA?



# Unpredictable social behaviour?

Remember the Turing halting problem?




In SPD, will a single strategy prove triumphant in the sense of progressively conquering more and more territory without opposition, or will an equilibrium of some small number of strategies emerge?

```
DEFINE DOES IT HALT (PROGRAM):  
{  
    RETURN TRUE;  
}
```

THE BIG PICTURE SOLUTION  
TO THE HALTING PROBLEM

Spatial PD systems are in general unpredictable...

# calvin and Hobbes



THEN I THOUGHT, LOOK, CHEATING ON ONE LITTLE TEST ISN'T SUCH A BIG DEAL. IT DOESN'T HURT ANYONE.

...BUT THEN I WONDERED IF I WAS JUST RATIONALIZING MY UNWILLINGNESS TO ACCEPT THE CONSEQUENCE OF NOT STUDYING.

